

# Assignment 1 - STATS 720

## Question 1: Olympic medals

Analyze the Olympic data set found [here](#) (raw download from [here](#)). The variables are:

- `team`: (approximately) country
  - `year`
  - `medal` (bronze/gold/silver)
  - `n`: medal count
  - `gdp`: GDP in const 2015 US\$ (billions)
  - `pop`: population size (millions)
- a. State which possible predictor variables you're going to include; justify your choice (refer to Harrell chapter 4 for rules of thumb about appropriate numbers of predictors).
- decide whether you're going to predict gold medals only, total medal count, or some weighted average of medals (e.g.  $4 \cdot G + 5 \cdot S + 2 \cdot B$ ).

## Load the data

```
library(RCurl)
library(readr)

urlfile <- "https://raw.githubusercontent.com/bbolker/stats720/main/data/olymp1.csv"

data<-read.csv(url(urlfile))

head(data)
```

	team	year	medal	n	gdp	pop
1	Afghanistan	2000	Bronze	0	6.206548	19.54298
2	Afghanistan	2000	Gold	0	6.206548	19.54298
3	Afghanistan	2000	Silver	0	6.206548	19.54298
4	Afghanistan	2004	Bronze	0	7.978516	23.55355
5	Afghanistan	2004	Gold	0	7.978516	23.55355
6	Afghanistan	2004	Silver	0	7.978516	23.55355

```
str(data)
```

```
'data.frame': 2435 obs. of 6 variables:
 $ team : chr "Afghanistan" "Afghanistan" "Afghanistan" "Afghanistan" ...
 $ year : int 2000 2000 2000 2004 2004 2004 2008 2008 2012 ...
 $ medal: chr "Bronze" "Gold" "Silver" "Bronze" ...
 $ n : int 0 0 0 0 0 1 0 0 1 ...
 $ gdp : num 6.21 6.21 6.21 7.98 7.98 ...
 $ pop : num 19.5 19.5 19.5 23.6 23.6 ...
```

```
#View(data)
```

**a. State which possible predictor variables you're going to include; justify your choice (refer to Harrell chapter 4 for rules of thumb about appropriate numbers of predictors).**

```
library(tidyverse)

wt_data<-data |>
  mutate(across(medal, ~ factor(., levels = c("Bronze", "Silver", "Gold")))) |>
  group_by(team, year) |>
  arrange(medal) |>
  summarise(n_wt = sum(c(1,2,4)*n)/4,
            gdp = mean(gdp),
            pop = mean(pop),
            .groups = "drop")|>
  na.omit()

#View(wt_data)
summary(wt_data)
```

team	year	n_wt	gdp
Length:541	Min. :2000	Min. : 0.00	Min. : 0.717
Class :character	1st Qu.:2004	1st Qu.: 0.00	1st Qu.: 21.592
Mode :character	Median :2008	Median : 1.50	Median : 87.348
	Mean :2008	Mean : 10.57	Mean : 553.276
	3rd Qu.:2012	3rd Qu.: 7.50	3rd Qu.: 331.035
	Max. :2016	Max. :202.00	Max. :18627.888


pop
Min. : 0.1074
1st Qu.: 3.9273
Median : 10.5148
Mean : 53.2954
3rd Qu.: 38.2586
Max. :1387.7900

### Descriptive statistics for predictors

```
library(ggplot2)
library(GGally)
```

```
Registered S3 method overwritten by 'GGally':
  method from
+.gg    ggplot2
```

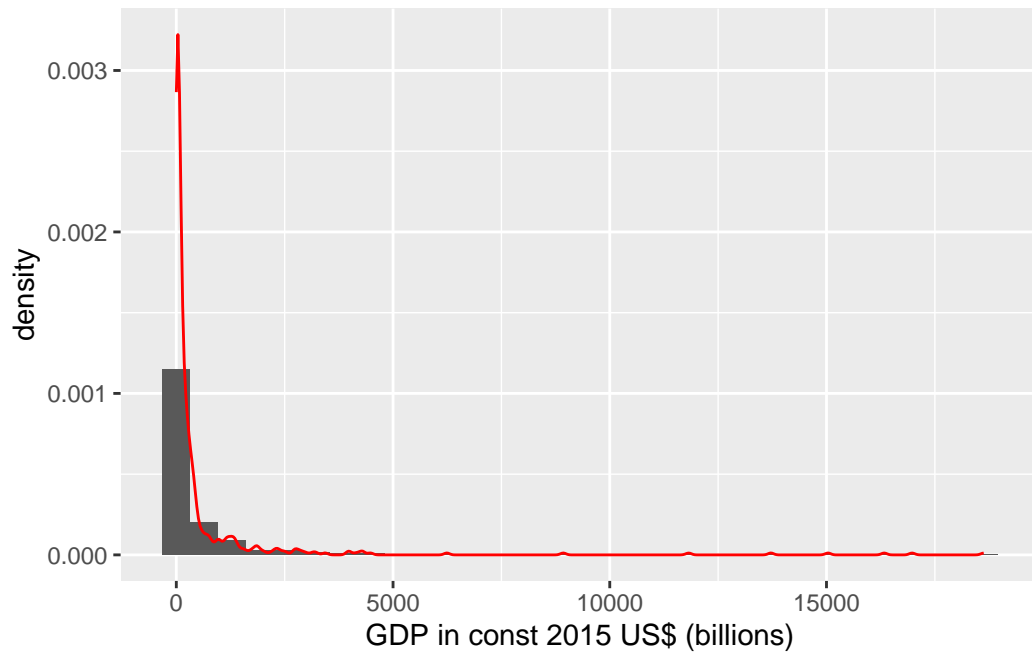
```
# Pairs plot between GDp and population
ggpairs(wt_data[,4:5])
```



```
#Histogram of GDP
ggplot(wt_data, aes(gdp)) +
  geom_histogram(aes(y=..density..)) + # scale histogram y
  geom_density(col = "red")+
  labs(x = "GDP in const 2015 US$ (billions)")
```


Warning: The dot-dot notation (`..density..`) was deprecated in ggplot2 3.4.0.  
i Please use `after\_stat(density)` instead.

`stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



```
# Histogram of Population
ggplot(wt_data, aes(pop)) +
  geom_histogram(aes(y=..density..)) + # scale histogram y    geom_density(col = "red")+
  labs(x = "Population size (millions)")
```

``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.



### Predictors:

To predict the medals, I am going to use Year, GDP, and population of the country. Natural spline with 5 df is used for GDP and Population. I also expect the interaction between GDP and year also population and Year. Since the GDP and population change over time. Total sample size is 541 which satisfy the Harrell rule of thumb of co-variate and sample size ratio 1:15

**Decide whether you're going to predict gold medals only, total medal count, or some weighted average of medals (e.g.  $4 \cdot G + 5 \cdot S + 2 \cdot B$ ). You can derive these different responses as follows:**

### Outcome:

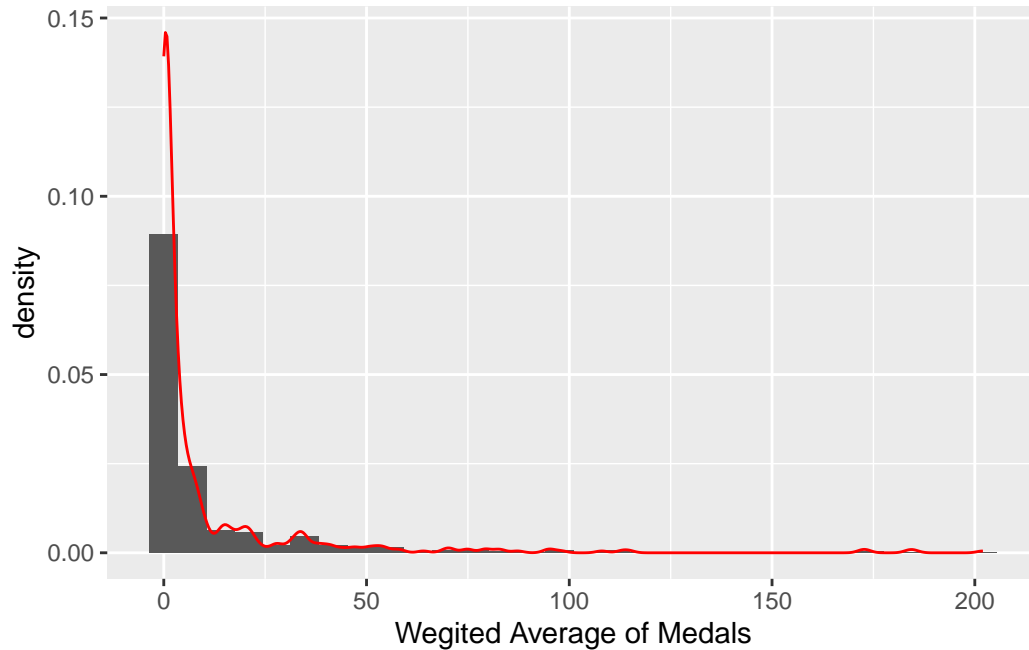
In this analysis, I use weighted average of medals as the outcome, because gold, silver, and bronze medals have different values of importance.

b. State the units of the response variable and of each predictor variable you plan to include; for each variable, state what you would consider as a reasonable threshold for a small change in that variable, *or* for a small slope (regression coefficient)

#### Histogram of Weighted average of medal

```
ggplot(wt_data, aes(n_wt)) +  
  geom_histogram(aes(y=..density..)) + # scale histogram y  
  geom_density(col = "red")+  
  labs(x = "Wegited Average of Medals")
```

`stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



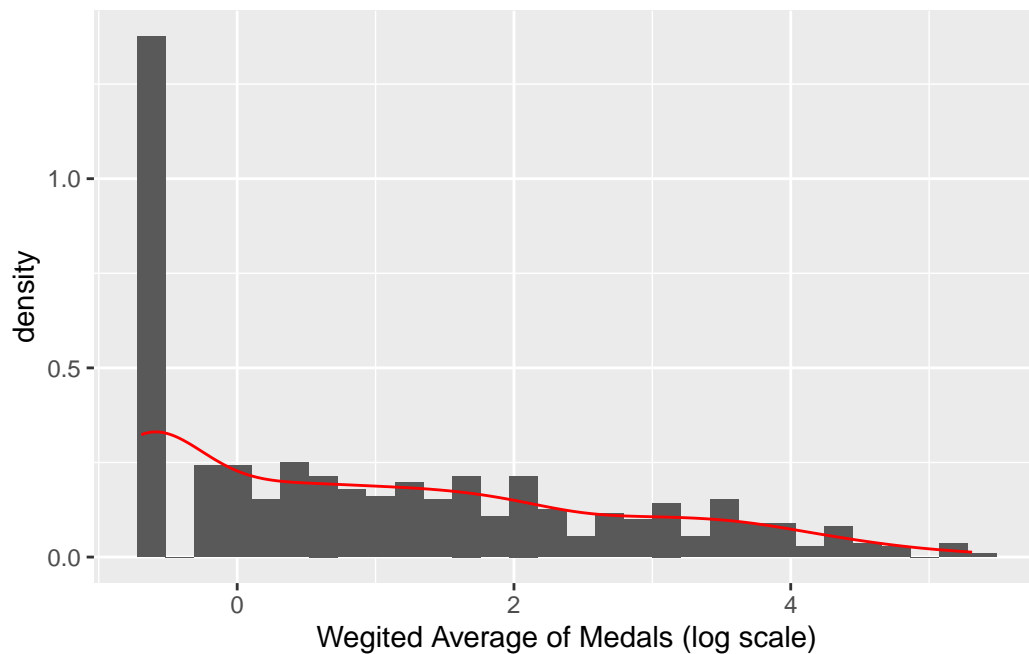
As the outcome is skewed, log transformation is used for in the linear model. And the outcome doesn't have any units.

- GDP is consider US\$ (billions) and
- Population size is consider in millions

```
wt_data$ln.n_wt<-log(wt_data$n_wt+0.5)

ggplot(wt_data, aes(ln.n_wt)) +
  geom_histogram(aes(y=..density..)) + # scale histogram y
  geom_density(col = "red")+
  labs(x = "Wegited Average of Medals (log scale)")
```

`stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



## Descriptive Statistics

```
cat("Summary of weighted number of medal:", "\n")
```

Summary of weighted number of medal:

```
summary(wt_data$n_wt)
```



Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.00	0.00	1.50	10.57	7.50	202.00

```
cat("Summary of weighted number of medal (log scale):","\n")
```

Summary of weighted number of medal (log scale):

```
summary(wt_data$ln.n_wt)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-0.6931	-0.6931	0.6931	1.0076	2.0794	5.3107

```
cat("Summary of GDP:", "\n")
```

Summary of GDP:

```
summary(wt_data$gdp)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.717	21.592	87.348	553.276	331.035	18627.888

```
cat("Summary of Population:", "\n")
```

Summary of Population:

```
summary(wt_data$pop)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.1074	3.9273	10.5148	53.2954	38.2586	1387.7900

## c. Fit Model

### Linear Model

```
library(splines)

Mod<-lm(ln.n_wt~year*ns(gdp, df=5)+year*ns(pop, df=5),wt_data)

summary(Mod)
```

Call:

```
lm(formula = ln.n_wt ~ year * ns(gdp, df = 5) + year * ns(pop,
  df = 5), data = wt_data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.4495	-0.8121	-0.1225	0.7688	2.9436

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-1.545e+02	7.800e+01	-1.981	0.0481 *
year	7.679e-02	3.886e-02	1.976	0.0487 *
ns(gdp, df = 5)1	-2.725e+01	9.696e+01	-0.281	0.7788
ns(gdp, df = 5)2	1.583e+02	9.045e+01	1.750	0.0807 .
ns(gdp, df = 5)3	-1.019e+02	2.665e+02	-0.382	0.7023
ns(gdp, df = 5)4	9.327e+01	2.332e+02	0.400	0.6894
ns(gdp, df = 5)5	-3.071e+02	2.930e+02	-1.048	0.2950
ns(pop, df = 5)1	1.659e+02	9.069e+01	1.829	0.0680 .
ns(pop, df = 5)2	7.663e+01	9.903e+01	0.774	0.4394
ns(pop, df = 5)3	1.473e+01	3.250e+02	0.045	0.9639
ns(pop, df = 5)4	1.993e+02	2.501e+02	0.797	0.4258
ns(pop, df = 5)5	3.476e+02	1.860e+02	1.869	0.0622 .
year:ns(gdp, df = 5)1	1.405e-02	4.830e-02	0.291	0.7713
year:ns(gdp, df = 5)2	-7.808e-02	4.506e-02	-1.733	0.0837 .
year:ns(gdp, df = 5)3	5.583e-02	1.327e-01	0.421	0.6741
year:ns(gdp, df = 5)4	-4.243e-02	1.161e-01	-0.365	0.7151
year:ns(gdp, df = 5)5	1.552e-01	1.457e-01	1.065	0.2875
year:ns(pop, df = 5)1	-8.246e-02	4.518e-02	-1.825	0.0685 .
year:ns(pop, df = 5)2	-3.812e-02	4.933e-02	-0.773	0.4400
year:ns(pop, df = 5)3	-7.749e-03	1.618e-01	-0.048	0.9618
year:ns(pop, df = 5)4	-1.000e-01	1.246e-01	-0.803	0.4224
year:ns(pop, df = 5)5	-1.742e-01	9.259e-02	-1.882	0.0604 .

---


Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

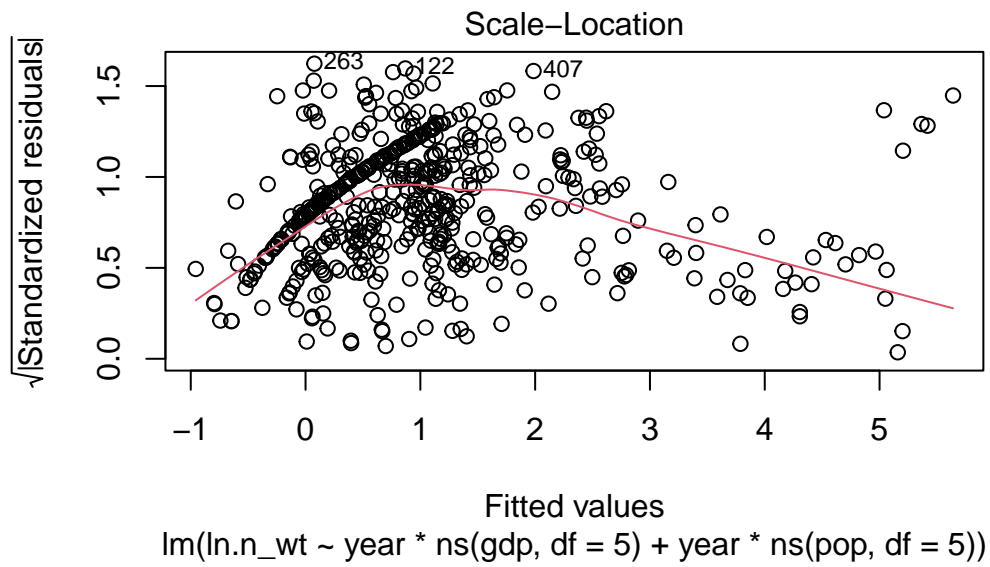

Residual standard error: 1.136 on 519 degrees of freedom  
Multiple R-squared: 0.5142, Adjusted R-squared: 0.4945  
F-statistic: 26.16 on 21 and 519 DF, p-value: < 2.2e-16


#### d. Diagnose the model

##### Performance plot

```
#par(mfrow=c(2,2))  
plot(Mod)
```







#### Interpretation:

##### Linearity Assumption:

As there is no clear pattern in the scatter plot of fitted values vs Residual. Therefore linearity assumption is satisfied.

##### Normality:

From the Q-Q plot, we can see that the points at the lower tail are slightly deviates from the diagonal line. However, in my opinion the normality assumption is satisfied approximately.

##### Homoscedasticity:

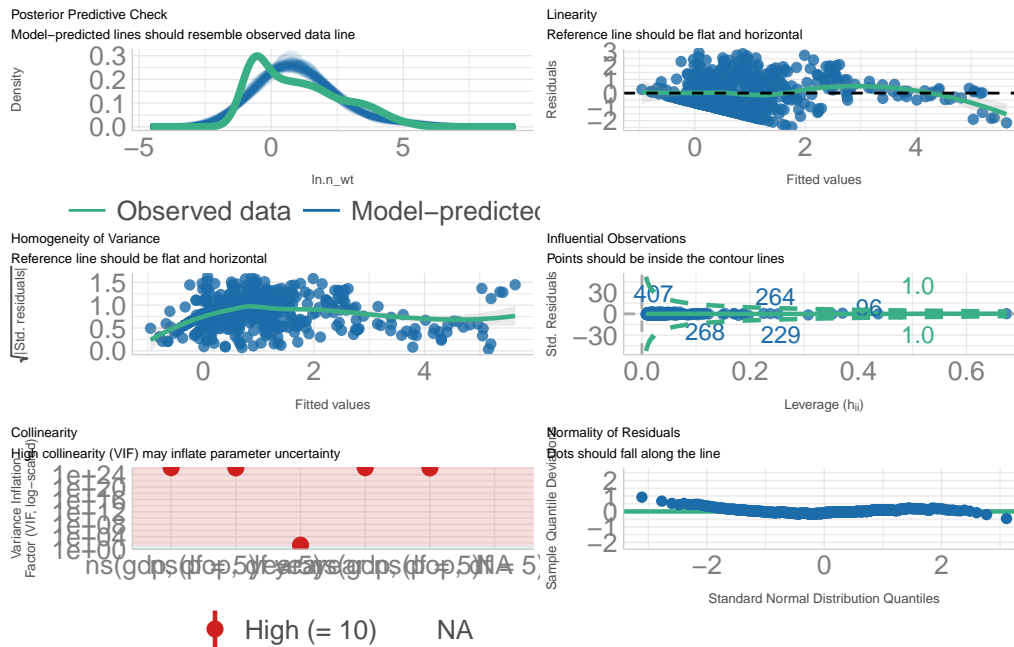
Scale-Location plot is dense upto 3 and then the spread is wider. Also, the red line is not horizontal until 3. It suggests that the homoscedasticity assumption is violated slightly.

##### Residual vs Leverage:

The plot depicts that there is no influential cases since all the points are inside the boundary.

```
library(performance)

check_model(Mod,panel=T,check='all',title_size=5,base_size=5,axis_title_size=5)
```



Performance plot is also says the same story of linearity, homogeneity of variance, normality of residuals, and influential points. In addition to that, The posterior predictive check plot illustrate observed and predicted density curve from the model which means, there is a slight deviation in the curve when it reaches the peak. VIF plot illustrates that there is a high collinearity among the predictors. There I drop both interaction term from the model.

## DHARMA


```
library(DHARMA)
```

This is DHARMA 0.4.6. For overview type `'?DHARMA'`. For recent changes, type `news(package = 'DHARMA')`.

```
simulationOutput <- simulateResiduals(fittedModel = Mod)

plot(simulationOutput)
```

## DHARMA residual



The Q-Q residual plot suggests that the model satisfies the normality of residual, homogeneity assumption and there is no influential points.

The residual vs predicted plot suggests that more residuals are in the lower tail of the distribution than we expect.

### e. If the model has any problems, make adjustments

To correct multi-collinearity, I drop the interaction term. In order to compare the most important predictor, I would scale the predictors and make unitless.

#### Scale

```
wt_data$s.gdp<-scale(wt_data$gdp,center = FALSE,scale = TRUE)

wt_data$s.pop<-scale(wt_data$pop,center = FALSE,scale = TRUE)
```

## Fitting the Model

```
Mod.Scaled<-lm(ln.n_wt~year+ns(s.gdp, df=5)+ns(s.pop, df=5),wt_data)

summary(Mod.Scaled)
```

Call:

```
lm(formula = ln.n_wt ~ year + ns(s.gdp, df = 5) + ns(s.pop, df = 5),
    data = wt_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.4189	-0.8536	-0.1652	0.7673	3.0655

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	28.130200	17.625388	1.596	0.11108
year	-0.014234	0.008781	-1.621	0.10559
ns(s.gdp, df = 5)1	1.007928	0.270050	3.732	0.00021 ***
ns(s.gdp, df = 5)2	1.585695	0.250029	6.342	4.87e-10 ***
ns(s.gdp, df = 5)3	9.952518	0.730633	13.622	< 2e-16 ***
ns(s.gdp, df = 5)4	8.312507	0.641493	12.958	< 2e-16 ***
ns(s.gdp, df = 5)5	4.603230	0.764216	6.023	3.20e-09 ***
ns(s.pop, df = 5)1	0.324111	0.254214	1.275	0.20288
ns(s.pop, df = 5)2	0.064812	0.277035	0.234	0.81512
ns(s.pop, df = 5)3	-0.558101	0.913552	-0.611	0.54152
ns(s.pop, df = 5)4	-1.391665	0.705981	-1.971	0.04922 *
ns(s.pop, df = 5)5	-2.191894	0.508931	-4.307	1.97e-05 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.145 on 529 degrees of freedom


Multiple R-squared: 0.4972, Adjusted R-squared: 0.4868



F-statistic: 47.56 on 11 and 529 DF, p-value: < 2.2e-16

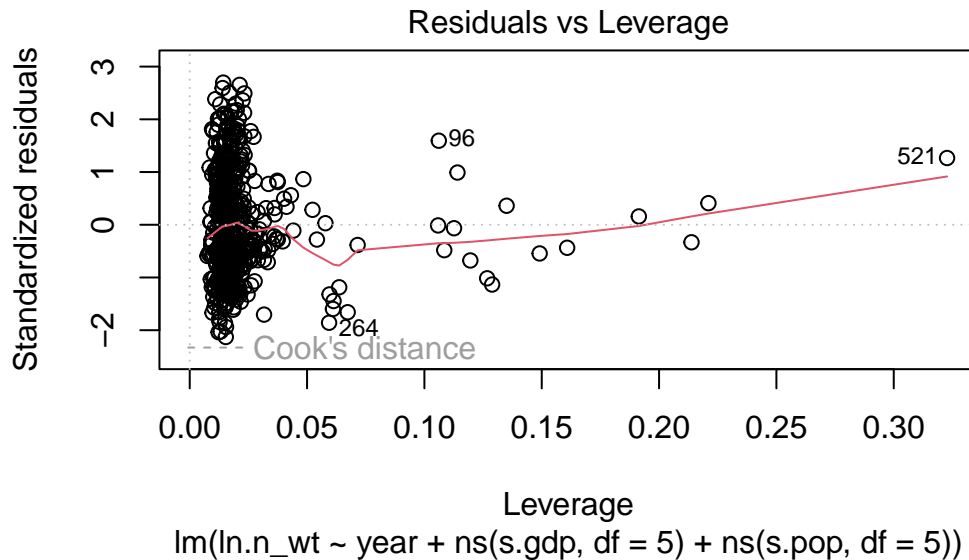


## Plots

```
plot(Mod.Scaled)
```







```
check_model(Mod.Scaled,panel=T,check='all',title_size=5,base_size=5,axis_title_size=5)
```

Some of the variables were in matrix-format - probably you used

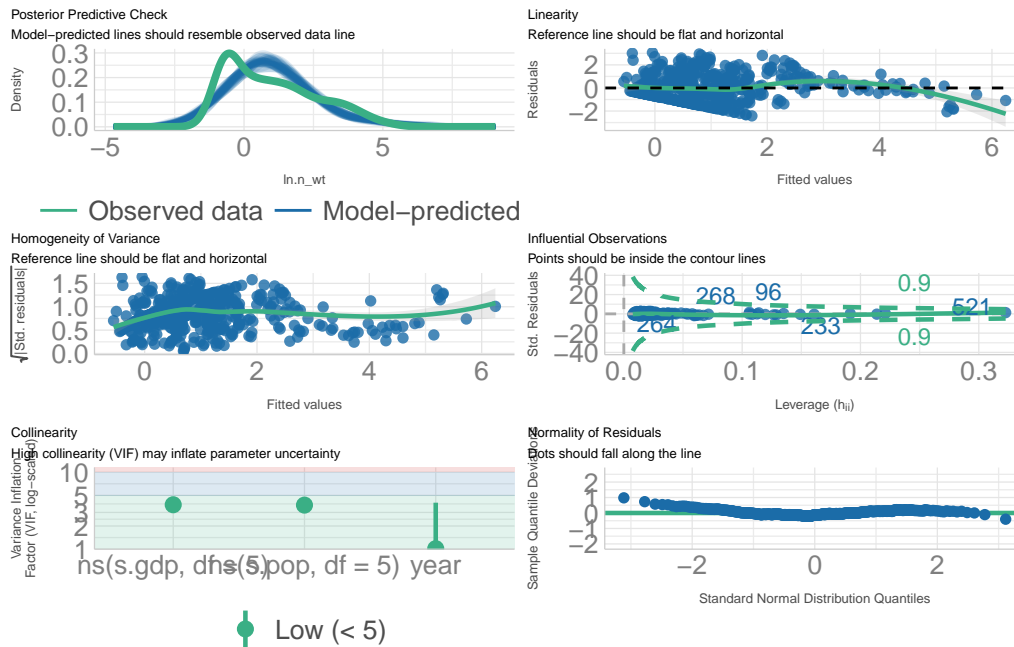
``scale()`` on your data?

If so, and you get an error, please try ``datawizard::standardize()`` to standardize your data.

Some of the variables were in matrix-format - probably you used

``scale()`` on your data?

If so, and you get an error, please try ``datawizard::standardize()`` to standardize your data.



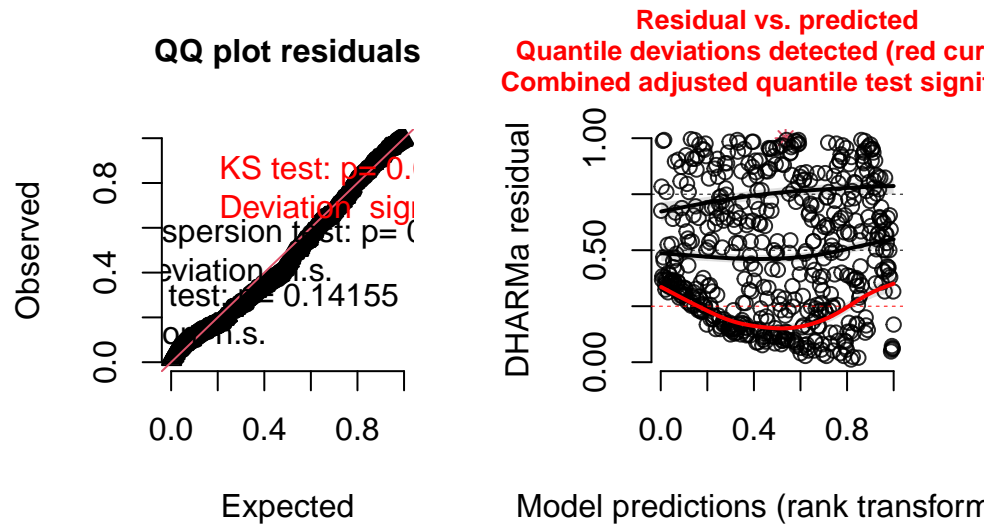
From the VIF plot, we can see that there is no collinear problem among predictors.

## DHARMA

```
simulationOutput.scaled <- simulateResiduals(fittedModel = Mod.Scaled)

plot(simulationOutput.scaled)
```

## DHARMA residual




The Q-Q residual plot suggests that the model satisfies the homogeneity assumption and there is no influential points. The residual vs predicted plot suggests that more residuals are in the lower tail of the distribution than we expect.

Since the sample size is large it can detect a small deviation of the normality. KS test is sensitive to sample size.

### f. Show a coefficient plot of the results

```
library(dotwhisker)

dwplot(Mod.Scaled)
```



From the plot we can see that GDP has a non-linear effect on the average number of medals. One year increase reduces the average number of medals by 1% ( $\exp(-0.01)$ ). As GDP and population are included as spline terms, the regression coefficients can't be interpreted directly.

#### g. Show an effects plot (predicted values or effects)

```
library(effects)
```

Loading required package: carData

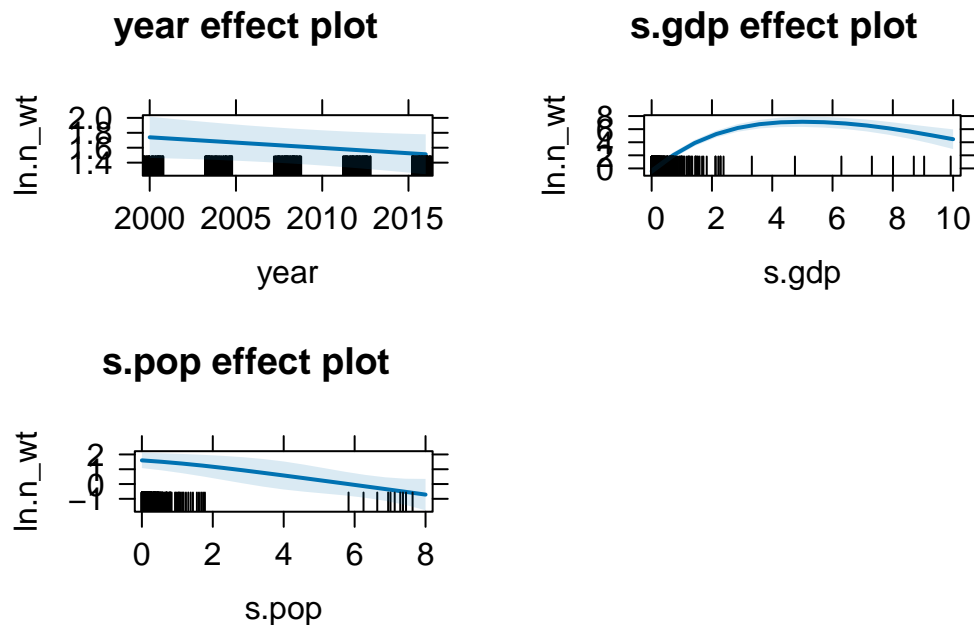
```
lattice theme set by effectsTheme()
See ?effectsTheme for details.
```

```
#effects::allEffects(Mod)

effect.plot<-effects::allEffects(Mod.Scaled)
```

```
Warning in Analyze.model(focal.predictors, mod, xlevels, default.levels, : the
predictors s.gdp, s.pop are one-column matrices that were converted to vectors
Warning in Analyze.model(focal.predictors, mod, xlevels, default.levels, : the
predictors s.gdp, s.pop are one-column matrices that were converted to vectors
Warning in Analyze.model(focal.predictors, mod, xlevels, default.levels, : the
predictors s.gdp, s.pop are one-column matrices that were converted to vectors
```

```
plot(effect.plot)
```



The marginal effect plots show that the year and population has negative effect on average medals, whereas GDP has quadratic effect.

### Estimate Marginal means

```
library(emmeans)
```

Welcome to emmeans.

Caution: You lose important information if you filter this package's results.  
See '? untidy'


Attaching package: 'emmeans'

The following object is masked from 'package:GGally':

pigs

```
plot(effect("year",Mod.Scaled))
```

Warning in Analyze.model(focal.predictors, mod, xlevels, default.levels, : the predictors s.gdp, s.pop are one-column matrices that were converted to vectors

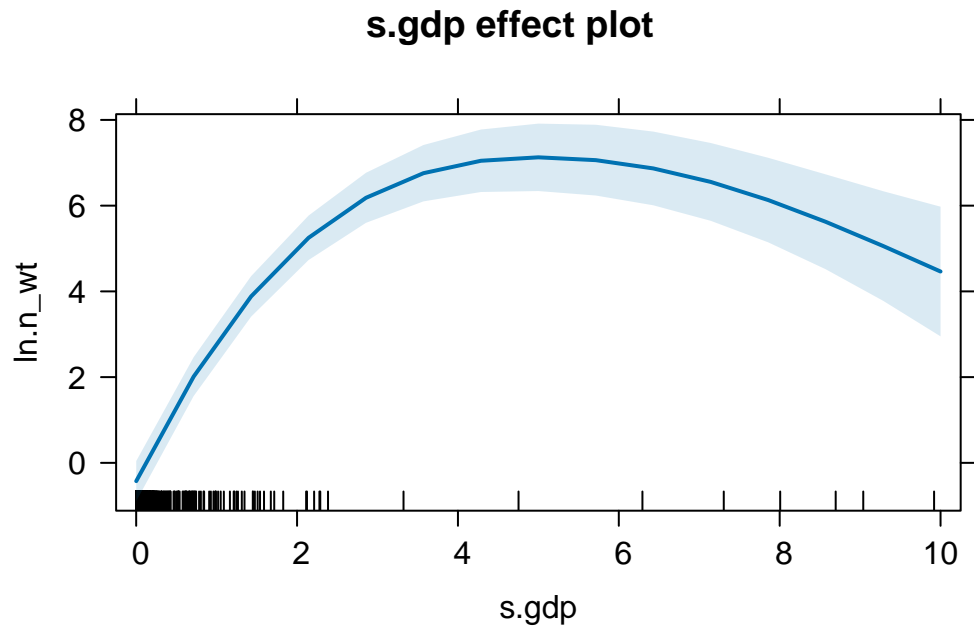


```
plot(effect("s.gdp",Mod.Scaled))
```

NOTE: s.gdp does not appear in the model

Warning in Analyze.model(focal.predictors, mod, xlevels, default.levels, : the predictors s.gdp, s.pop are one-column matrices that were converted to vectors






```
plot(effect("s.pop",Mod.Scaled))
```

NOTE: s.pop does not appear in the model

Warning in Analyze.model(focal.predictors, mod, xlevels, default.levels, : the predictors s.gdp, s.pop are one-column matrices that were converted to vectors



Test for non-linearity

GDP

$$H_0 : \beta_3 = \beta_4 = \beta_5 = \beta_6$$

```
library(car)
```

Attaching package: 'car'

The following object is masked from 'package:dplyr':

recode

The following object is masked from 'package:purrr':

some

```
linearHypothesis(Mod, names(coef(Mod))[4:7])
```

Linear hypothesis test

Hypothesis:

`ns(gdp, df = 5)2 = 0`

`ns(gdp, df = 5)3 = 0`

`ns(gdp, df = 5)4 = 0`

`ns(gdp, df = 5)5 = 0`

Model 1: restricted model

Model 2: `ln.n_wt ~ year * ns(gdp, df = 5) + year * ns(pop, df = 5)`

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	523	679.33				
2	519	670.13	4	9.1992	1.7811	0.1312

ANOVA/F test suggests that the full model is has significant lower RSS, means natural spline of GDP has non-linear effect on log transformed weighted average medal.

**Population**

$$H_0 : \beta_9 = \beta_{10} = \beta_{11} = \beta_{12}$$

```
linearHypothesis(Mod, names(coef(Mod))[9:12])
```

Linear hypothesis test

Hypothesis:

`ns(pop, df = 5)2 = 0`

`ns(pop, df = 5)3 = 0`

`ns(pop, df = 5)4 = 0`

`ns(pop, df = 5)5 = 0`

Model 1: restricted model

Model 2: `ln.n_wt ~ year * ns(gdp, df = 5) + year * ns(pop, df = 5)`

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	523	675.19				
2	519	670.13	4	5.0537	0.9785	0.4188

ANOVA/F test suggests that the full model is has significant lower RSS, means natural spline of population has non-linear effect on log transformed weighted average medal.

## Question 2: contrasts

Suppose we have an experiment with four levels: control (C) and three increasing levels of the treatment (I, II, III). We are interested in:

- the difference between the control and the *average* of the treatment levels
- *successive differences* (I vs II, II vs III) among the non-control treatments.

Construct a set of contrasts to quantify these effects. Test your results by making up a minimal data frame with just one observation per treatment. Fit the linear model and show that the coefficients match what you intended

```
trt<-c("control","I","II","III")
mu<-c(0.5,0.9,1.2,0.7)

set.seed(100)
y<-rnorm(length(trt),mu,0.5)

da<-data.frame(y=y,x=trt)
da
```

```
      y      x
1 0.2489038 control
2 0.9657656      I
3 1.1605415     II
4 1.1433924     III
```

```
mod<-lm(y~as.factor(x),da)
mod
```

Call:

```
lm(formula = y ~ as.factor(x), data = da)
```

Coefficients:

(Intercept)	as.factor(x)I	as.factor(x)II	as.factor(x)III
0.2489	0.7169	0.9116	0.8945

```
Mean.eff<-predict(mod,da)
```

```
# 3 contrasts:
```

```
# 1. the difference between the control and the average of the treatment levels
```

```
# 2. difference of I vs II
```

```
# 3. difference II vs III among the non-control treatments
```

```
C.inv<-matrix(c(1,-1/3,-1/3,-1/3,  
               0,1,-1,0,  
               0,0,1,-1),  
             byrow = T,nrow = 3)
```

```
C.inv
```

```
      [,1]      [,2]      [,3]      [,4]  
[1,]    1 -0.3333333 -0.3333333 -0.3333333  
[2,]    0  1.0000000 -1.0000000  0.0000000  
[3,]    0  0.0000000  1.0000000 -1.0000000
```

```
C.inv%%Mean.eff
```

```
      [,1]  
[1,] -0.84099599  
[2,] -0.19477587  
[3,]  0.01714905
```

```
# C matrix with inercept
```

```
C<-matrix(c(1,1,1,1,  
            1,-1/3,-1/3,-1/3,  
            0,1,-1,0,  
            0,0,1,-1),  
          byrow = F,nrow = 4)
```

```
cat("C matrix","\n")
```

```
C matrix
```

```
C
```

```
      [,1]      [,2] [,3] [,4]
[1,]      1 1.0000000      0      0
[2,]      1 -0.3333333      1      0
[3,]      1 -0.3333333     -1      1
[4,]      1 -0.3333333      0     -1
```

```
library(MASS)
```

```
Attaching package: 'MASS'
```

```
The following object is masked from 'package:dplyr':
```

```
select
```

```
# C-inverse matrix
```

```
cat("C-inverse matrix","\n")
```

```
C-inverse matrix
```

```
Contrast.mat<-fractions(solve(C))
Contrast.mat[,-1]
```

```
      [,1] [,2] [,3]
[1,]  1/4  1/4  1/4
[2,] -1/4 -1/4 -1/4
[3,]  2/3 -1/3 -1/3
[4,]  1/3  1/3 -2/3
```

### Question 3: simulations to evaluate the effects of model misspecification

#### Function to Simulate data from t-distribution

```

sim_fun <- function(n = 100, slope = 1, sd = 1, intercept = 0,df=2) {
  x <- runif(n)
  mu<-intercept + slope * x
  #y <- rnorm(n, intercept + slope * x, sd = sd)
  y<-mu+sd*rt(n, df)
  data.frame(x, y)
}

#sim<-sim_fun (n = 100, slope = 1, sd = 1, intercept = 0)

```

### Evaluate coverage for model mis-specification

```

Model_misp<-function(n=100,t.df=2,slope=1,sd=1,intercept=0,B=1000,alpha=0.05)
{

  out<-data.frame(matrix(0,nrow=B,ncol = 3))
  colnames(out)<-c("slope","p.val","Coverage")

  for(i in 1:B)
  {
    #cat("Iteration number is:",i,"\n")

    sim.dat<-sim_fun (n = n, slope = slope, sd = sd, intercept = intercept,df=t.df)

    head(sim.dat)

    lm.mod<-lm(y~x,sim.dat)
    out$slope[i] <- coef(lm.mod)[2]
    out$p.val[i]<-coef(summary(lm.mod))[2, "Pr(>|t|)"]

    between <- function(a, b) (b[1] < a & a < b[2]);
    out$Coverage[i]<- between(slope, confint(lm.mod)[2,])
  }

  Bias<-mean(out$slope-slope)
  SE<-sd(out$slope)
  RMSE<-sqrt(mean((out$slope-slope)^2))

  Power<-mean(out$p.val<alpha)
}

```

```

Coverage<-mean(out$Coverage)

#Metrics<-list(Bias=Bias,SE=SE,RMSE=RMSE,Power=Power,Coverage=Coverage)

return(results=list(Bias=Bias,SE=SE,RMSE=RMSE,Power=Power,Coverage=Coverage,Estimates=ou
})

res<-Model_misp(n=100,t.df=2,slope=1,sd=1,intercept=0,B=1000,alpha=0.05)

n<-seq(10,100,by=10)
df<-seq(2, 50, by = 6)

results<-expand.grid(n=n,df=df)

for(i in 1:nrow(results)){

  set.seed(1000+i)
  res<-Model_misp(n=results$n[i],t.df=results$df[i],slope=1,sd=1,intercept=0,B=1000,alp

results$Bias[i]<-res$Bias
results$SE[i]<-res$SE
results$RMSE[i]<-res$RMSE
results$Power[i]<-res$Power
results$Coverage[i]<-res$Coverage

}

results

```

	n	df	Bias	SE	RMSE	Power	Coverage
1	10	2	0.0218595214	3.5895902	3.5878615	0.076	0.937
2	20	2	-0.0073709893	2.6444793	2.6431670	0.105	0.958
3	30	2	-0.0778527608	2.0815568	2.0819719	0.116	0.958
4	40	2	-0.1017981934	1.8502557	1.8521300	0.147	0.940
5	50	2	-0.0347609093	1.4947178	1.4943746	0.145	0.954
6	60	2	0.0968219980	1.8208437	1.8225067	0.192	0.956
7	70	2	0.0024580897	1.1240820	1.1235225	0.199	0.961
8	80	2	0.0237283434	1.4402335	1.4397088	0.223	0.956
9	90	2	-0.0002820687	1.2096921	1.2090871	0.223	0.953
10	100	2	-0.0126257102	1.1320592	1.1315635	0.250	0.951



11	10	8	-0.0515679820	1.4259321	1.4261515	0.110	0.955
12	20	8	-0.0067655281	0.9552187	0.9547649	0.196	0.942
13	30	8	-0.0151260646	0.7546581	0.7544323	0.254	0.949
14	40	8	0.0125071126	0.6630010	0.6627874	0.349	0.944
15	50	8	0.0068930966	0.5953288	0.5950710	0.408	0.940
16	60	8	0.0204083386	0.5302600	0.5303876	0.495	0.952
17	70	8	0.0224191729	0.4769822	0.4772704	0.550	0.953
18	80	8	0.0122020768	0.4629180	0.4628474	0.616	0.944
19	90	8	0.0157803038	0.4466383	0.4466937	0.655	0.946
20	100	8	-0.0042286131	0.3908542	0.3906816	0.684	0.959
21	10	14	-0.0518997522	1.2932170	1.2936118	0.097	0.951
22	20	14	-0.0058549002	0.8709583	0.8705424	0.198	0.944
23	30	14	0.0060676230	0.7146047	0.7142731	0.270	0.949
24	40	14	-0.0186884651	0.6068555	0.6068398	0.362	0.950
25	50	14	0.0093959287	0.5514083	0.5512126	0.447	0.952
26	60	14	0.0133625731	0.4939028	0.4938366	0.522	0.950
27	70	14	-0.0071733354	0.4473936	0.4472274	0.588	0.953
28	80	14	-0.0104992466	0.4275915	0.4275066	0.618	0.949
29	90	14	-0.0156276852	0.4031473	0.4032486	0.695	0.941
30	100	14	-0.0021582989	0.3623533	0.3621785	0.780	0.947
31	10	20	-0.0568342173	1.3381028	1.3386407	0.117	0.935
32	20	20	-0.0279860007	0.8242102	0.8242732	0.199	0.961
33	30	20	-0.0122992554	0.6808521	0.6806227	0.297	0.952
34	40	20	-0.0091135232	0.5890949	0.5888708	0.399	0.961
35	50	20	0.0161912074	0.5185226	0.5185161	0.477	0.960
36	60	20	-0.0095241911	0.4772801	0.4771365	0.535	0.955
37	70	20	-0.0342428118	0.4665877	0.4676098	0.591	0.939
38	80	20	0.0037466861	0.4134585	0.4132687	0.668	0.952
39	90	20	-0.0132804259	0.4013590	0.4013781	0.694	0.944
40	100	20	-0.0021710612	0.3669792	0.3668021	0.770	0.950
41	10	26	-0.0187496641	1.2642145	1.2637213	0.128	0.949
42	20	26	0.0124275774	0.8517176	0.8513823	0.203	0.939
43	30	26	0.0006904761	0.6807617	0.6804216	0.320	0.945
44	40	26	-0.0081542237	0.5873802	0.5871431	0.406	0.946
45	50	26	0.0192417729	0.5073518	0.5074630	0.491	0.965
46	60	26	-0.0064943529	0.4851123	0.4849132	0.562	0.934
47	70	26	-0.0213907719	0.4534566	0.4537343	0.606	0.949
48	80	26	-0.0158201673	0.3939339	0.3940546	0.681	0.959
49	90	26	0.0073036795	0.3824970	0.3823754	0.735	0.956
50	100	26	-0.0104583120	0.3568978	0.3568725	0.778	0.954
51	10	32	-0.0241962664	1.2978057	1.2973823	0.118	0.937
52	20	32	-0.0031347924	0.8392026	0.8387887	0.209	0.946
53	30	32	0.0071653371	0.6923543	0.6920451	0.325	0.946

54	40	32	0.0189375934	0.5769600	0.5769823	0.413	0.953
55	50	32	-0.0343354944	0.5160919	0.5169753	0.462	0.946
56	60	32	0.0011451290	0.4690567	0.4688235	0.564	0.953
57	70	32	-0.0130525207	0.4376707	0.4376465	0.622	0.946
58	80	32	0.0012771787	0.4012028	0.4010042	0.695	0.950
59	90	32	0.0113366553	0.3866452	0.3866180	0.739	0.936
60	100	32	-0.0105588159	0.3662593	0.3662284	0.766	0.949
61	10	38	0.0370772057	1.2700155	1.2699217	0.128	0.935
62	20	38	0.0132638985	0.8245994	0.8242937	0.221	0.954
63	30	38	0.0288349296	0.6487288	0.6490452	0.306	0.954
64	40	38	-0.0149106178	0.5799881	0.5798898	0.404	0.954
65	50	38	-0.0103223888	0.4728997	0.4727758	0.493	0.967
66	60	38	0.0066717373	0.4693540	0.4691667	0.559	0.945
67	70	38	-0.0264588717	0.4250136	0.4256242	0.604	0.958
68	80	38	-0.0276799035	0.4077135	0.4084486	0.661	0.955
69	90	38	-0.0172105351	0.3859669	0.3861576	0.736	0.953
70	100	38	-0.0088593634	0.3452754	0.3452164	0.787	0.949
71	10	44	0.0123797648	1.2827362	1.2821544	0.123	0.946
72	20	44	0.0293393573	0.8272464	0.8273531	0.222	0.951
73	30	44	-0.0021412368	0.6462148	0.6458952	0.310	0.964
74	40	44	0.0107049715	0.5673222	0.5671395	0.422	0.952
75	50	44	-0.0113795459	0.5254795	0.5253399	0.487	0.943
76	60	44	-0.0165485090	0.4600446	0.4601123	0.553	0.956
77	70	44	0.0108335607	0.4288320	0.4287544	0.652	0.953
78	80	44	-0.0089833849	0.4035550	0.4034532	0.705	0.944
79	90	44	0.0097965436	0.3738644	0.3738058	0.750	0.949
80	100	44	-0.0090846682	0.3606650	0.3605991	0.786	0.941
81	10	50	0.0111526674	1.1860166	1.1854759	0.098	0.970
82	20	50	-0.0369001075	0.8618221	0.8621811	0.229	0.943
83	30	50	0.0191935240	0.6680284	0.6679701	0.322	0.961
84	40	50	-0.0076266623	0.5715790	0.5713441	0.416	0.949
85	50	50	-0.0283692453	0.5309160	0.5314083	0.483	0.937
86	60	50	-0.0287830245	0.4701709	0.4708164	0.544	0.940
87	70	50	-0.0222349591	0.4392776	0.4396206	0.606	0.951
88	80	50	0.0013181087	0.3987025	0.3985053	0.686	0.954
89	90	50	-0.0081176928	0.3758140	0.3757137	0.740	0.959
90	100	50	0.0020334640	0.3599674	0.3597931	0.792	0.948


```
#View(results)
```

## Plot - Results

### Bias

```
#cols <- c("gray", "gray", "gray","gray", "gray", "gray","gray", "gray", "blue")

ggplot(results, aes(x=n, y=Bias)) +
  geom_line(aes(color=as.factor(df)),linewidth =1)+
  guides(color = guide_legend(title = "DF")) #+
```



```
#scale_color_manual(values = cols)
```


Bias plot suggests that as sample size increases the bias converges to zero irrespective of the df.

### RMSE

```
#cols <- c("gray", "gray", "gray","gray", "gray", "gray","gray", "gray", "blue")

ggplot(results, aes(x=n, y=RMSE)) +
  geom_line(aes(color=as.factor(df)),linewidth =1)+
```

```
guides(color = guide_legend(title = "DF"))#+
```




```
#scale_color_manual(values = cols)
```

RMSE depicts that except  $df=2$ , RMSE are almost equal irrespective of the  $df$ .

## Power

```
cols <- c("gray", "gray", "gray", "gray", "gray", "gray", "gray", "gray", "blue")

ggplot(results, aes(x=n, y=Power)) +
  geom_line(aes(color=as.factor(df)), linewidth = 1) +
  guides(color = guide_legend(title = "DF"))
```




The plot suggest that as sample size and df increases the power increases.

### Coverage

```
cols <- c("gray", "gray", "gray", "gray", "gray", "gray", "gray", "gray", "blue")

ggplot(results, aes(x=n, y=Coverage)) +
  geom_line(aes(color=as.factor(df)), linewidth = 1) +
  guides(color = guide_legend(title = "DF"))
```



The coverage is highly oscillating irrespective of df.